

ECE 484 Digital Image Processing

Lec 05 - Linear Filtering

Zhu Li

Dept of CSEE, UMKC

Office: FH560E, Email: lizhu@umkc.edu, Ph: x 2346.

<http://l.web.umkc.edu/lizhu>



Outline

- Recap of Lec 04
- Linear Filters
- Smoothing and Edge Detection
- Accelerating Linear Filters
- Summary

Gamma Correction - Adjust dynamic ranges

□ Matching Display characteristics

$$v = c \cdot u^\gamma$$

- power-law response functions in practice

CRT Intensity-to-voltage function has
 $\gamma \approx 1.8\sim 2.5$

Camera capturing distortion with $\gamma_c = 1.0\sim 1.7$

Similar device curves in scanners, printers, ...

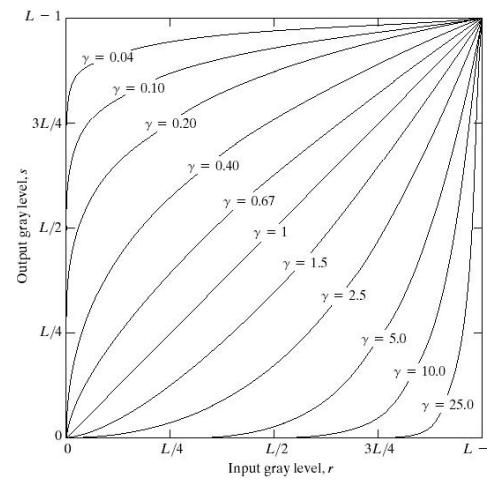


FIGURE 3.6 Plots of the equation $s = cr^\gamma$ for various values of γ ($c = 1$ in all cases).

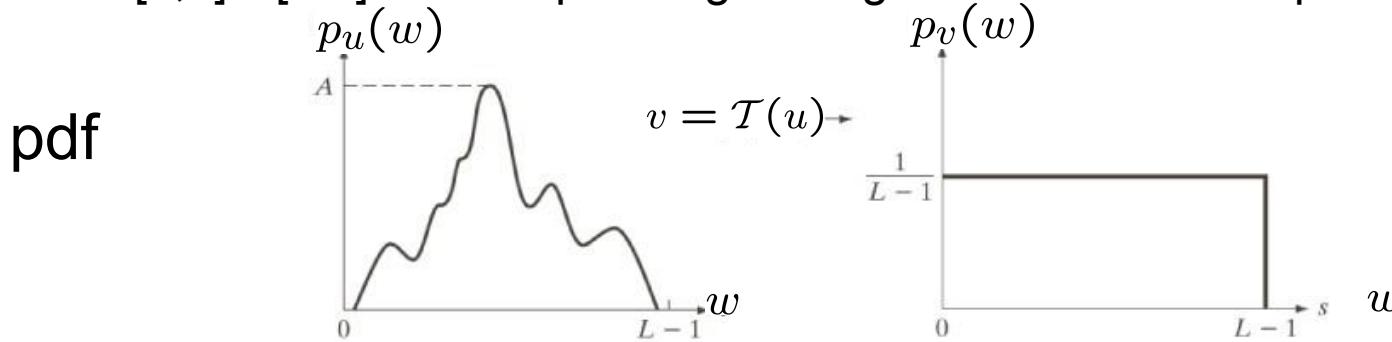


- power-law transformations are also useful for general purpose contrast manipulation

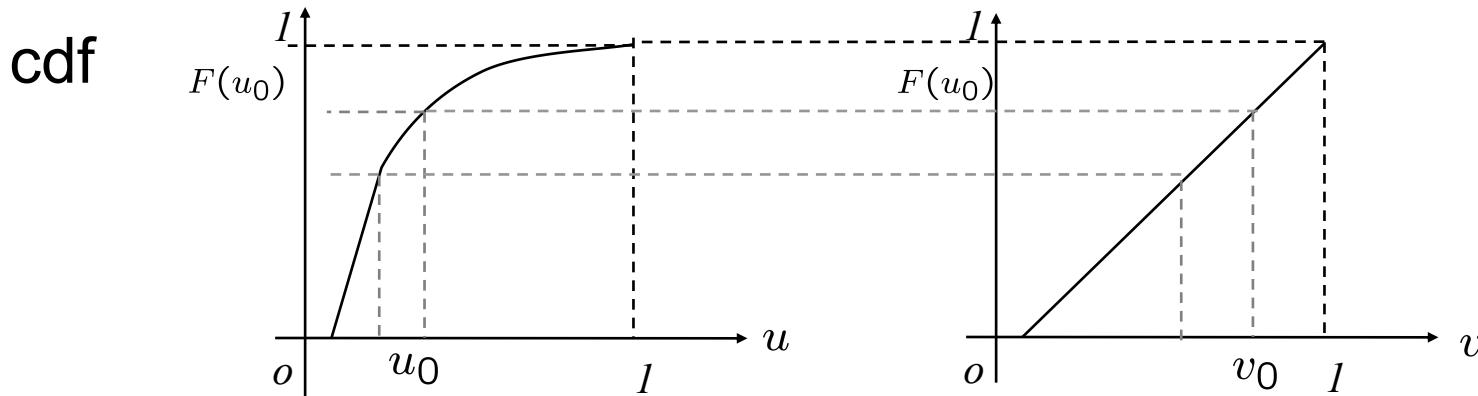
Histogram Equalization

❑ Objectives

- goal: map the each luminance level to a new value such that the output image has approximately uniform distribution of gray levels
- two desired properties
 - monotonic (non-decreasing) function: no value reversals
 - $[0,1] \rightarrow [0,1]$: the output range being the same as the input range



$$F_u(U) = P(U < u) = \int_0^u p_u(w) dw \quad F_v(V) = P(V < v) = \int_0^v p_v(w) dw$$



Histogram Equalization

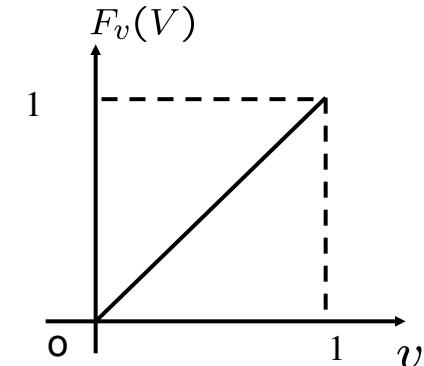
□ Algorithm

- make

$$v = F_u(u) = P(U < u) = \int_0^u p_u(w) dw$$

- show

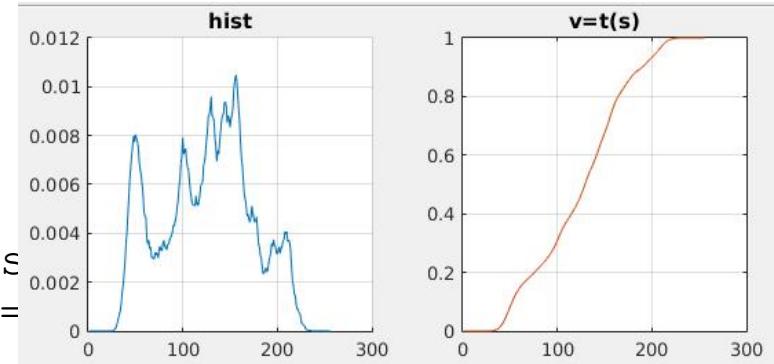
$$F_v(V) = v, \quad v \in [0, 1]$$



$$\begin{aligned} F_v(V) &= P(V < v) = P(F_u(U) < v) \\ &= P(U < F_u^{-1}(v)) = F_u(F_u^{-1}(v)) = v \end{aligned}$$

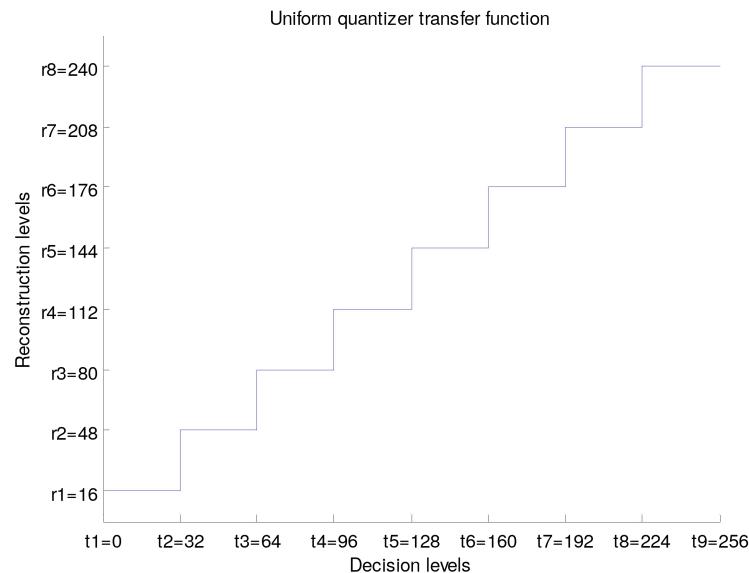
□ Matlab

```
im = imread('lena.png');
h1 = imhist(rgb2gray(im));
h1=h1/sum(h1);
cf1 = cumsum(h1);
subplot(1,2,1); plot(h1); title('his
subplot(1,2,2); plot(cf1); title('v=
```



Uniform Quantization: B=3, L=8

- ☐ Finer quantization: $B=3 \Rightarrow L=8$; false contours present



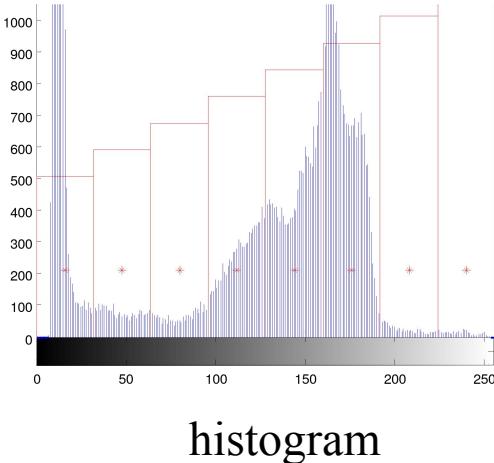
Non-quantized image



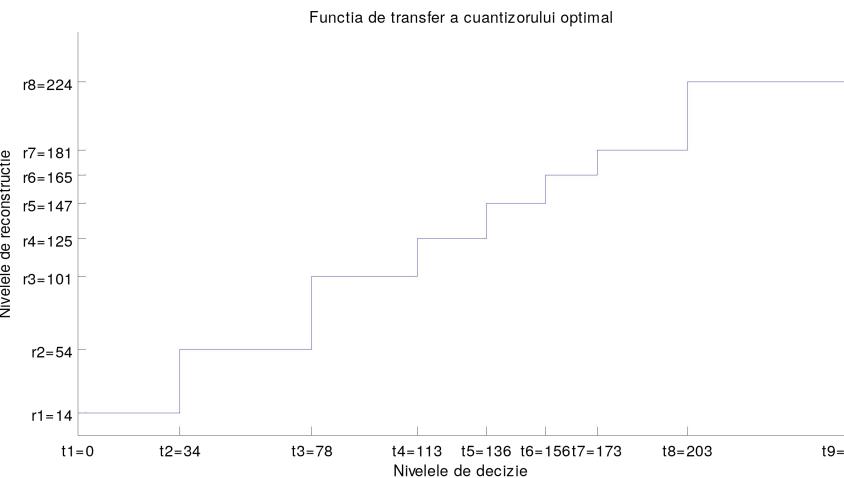
Quantized image



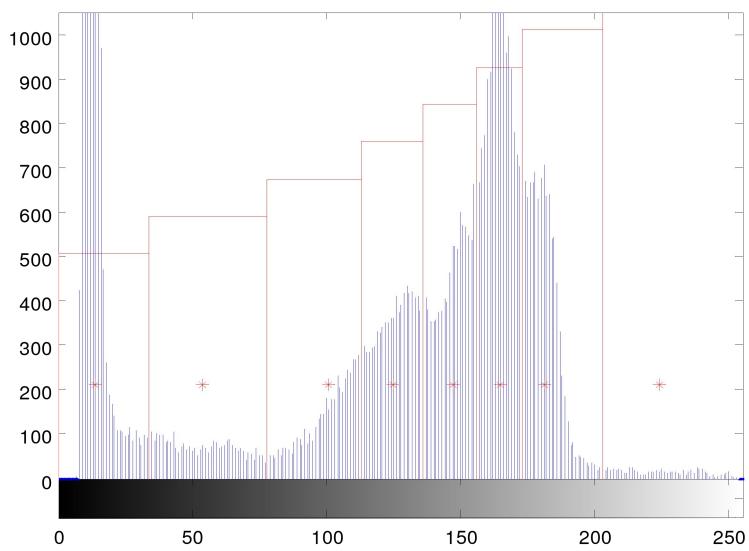
Quantization error; MSE=7.33



□ LMQ: B=3, L=8



Non-quantized image



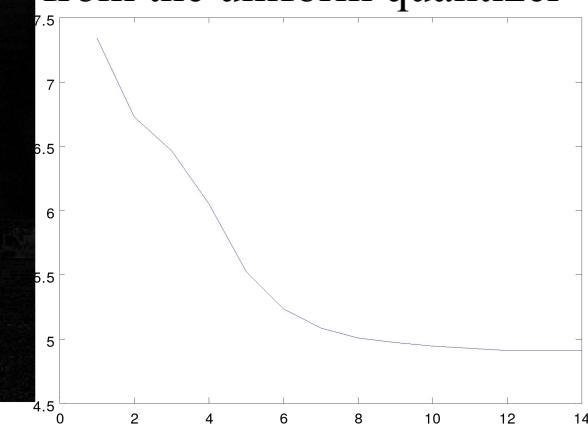
The quantization error; The evolution of MSE
MSE=5



Quantized image



The quantization error; The evolution of MSE
in the optimization, starting from the uniform quantizer

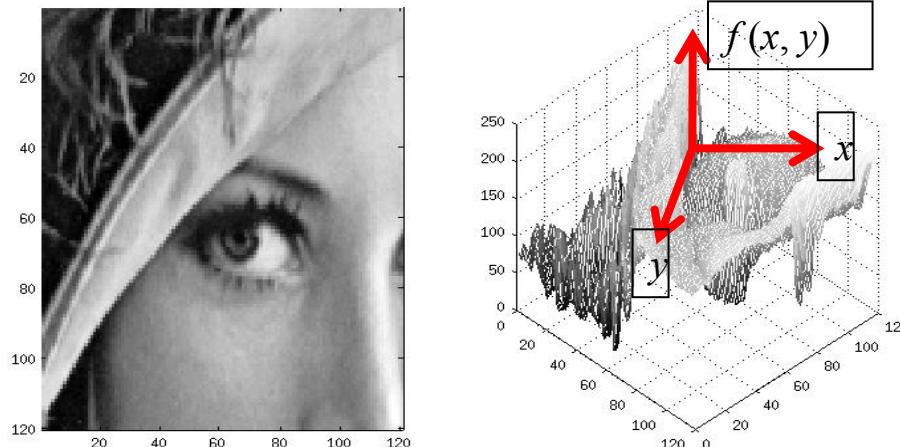


Outline

- Recap of Lec 04
- Linear Filters
- Smoothing and Edge Detection
- Accelerating Linear Filters
- Summary

What is an image?

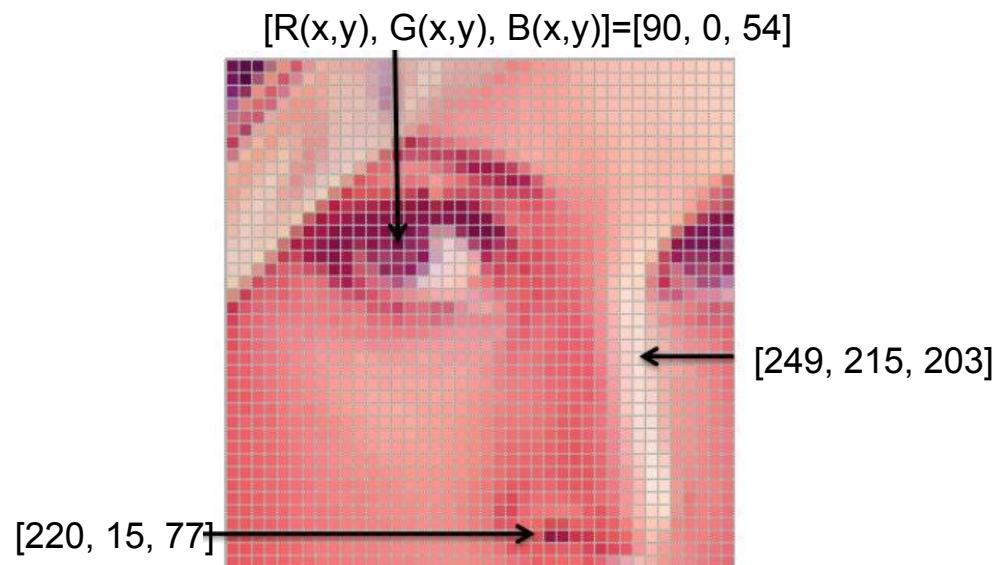
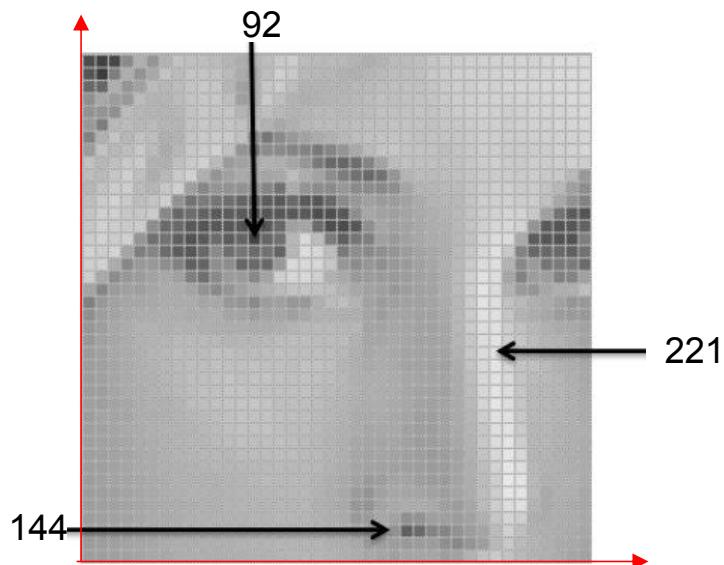
- We can think of a (grayscale) image as a **function**, f , from \mathbb{R}^2 to \mathbb{R} (or a 2D *signal*):
 - $f(x,y)$ gives the **intensity** at position (x,y)



- A **digital** image is a discrete (**sampled, quantized**) version of this function

Image as a function: $I=f(x,y)$;

- A grid (matrix) of intensity values



(common to use one byte per value: 0 = black, 255 = white)

Motivation: Why Filtering ?

□ Some use cases...

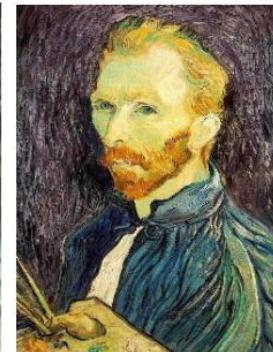
De-noising



Salt and pepper noise



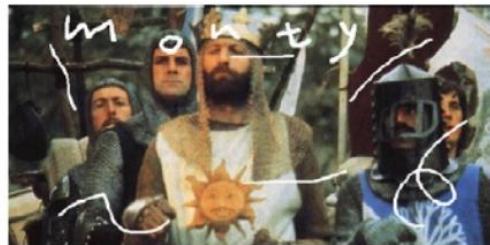
Super-resolution



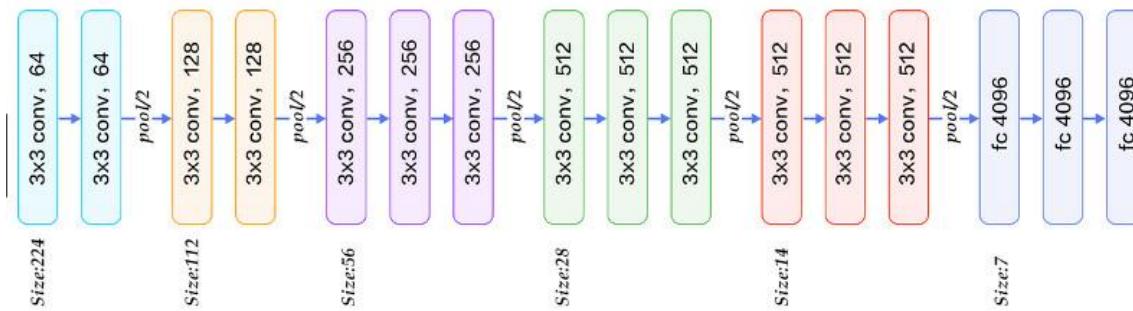
Edge detection



In-painting



Bertamio et al



Deep Learning with Conv Nets:
VGG16

Image filtering – Looking at pixel neighbors

- Modify the pixels in an image based on some function of a local neighborhood of each pixel
 - Discrete form:

$$f[n, m] \rightarrow \boxed{\text{System } \mathcal{S}} \rightarrow g[n, m]$$

$$g = \mathcal{S}[f], \quad g[n, m] = \mathcal{S}\{f[n, m]\}$$

$$f[n, m] \xrightarrow{\mathcal{S}} g[n, m]$$

10	5	3
4	5	1
1	1	7

Local image data

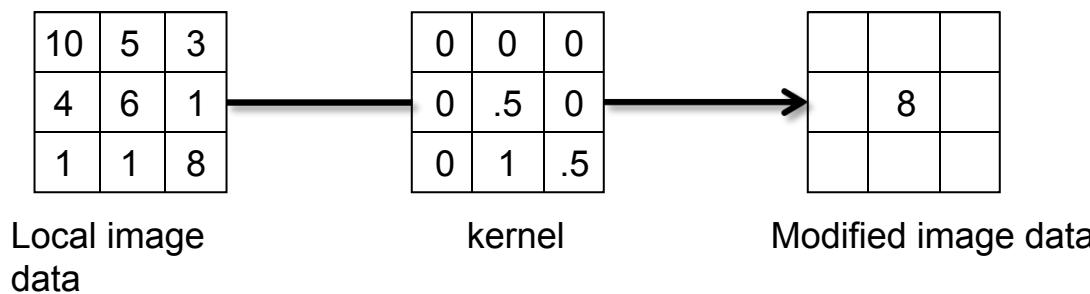


		7

Modified image data

Linear filtering

- One simple version: linear filtering (cross-correlation, convolution)
 - Replace each pixel by a linear combination of its neighbors
- The prescription for the linear combination is called the “*kernel*” (or “mask”, “filter”)



Source: L. Zhang

Cross-correlation

□ Cross-Correlation:

- No flipping of kernel

Let F be the image, H be the kernel
(of size $2k+1 \times 2k+1$), and G be the
output image

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v]F[i + u, j + v]$$

This is called a **cross-correlation**
operation:

$$G = H \otimes F$$

Convolution

- Same as cross-correlation, except that the kernel is “flipped” (horizontally and vertically), called convolution

$$G = H * F$$

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v]F[i - u, j - v]$$

- Convolution is **commutative** and **associative**
 - $F = G * h + G * f = G * (h + f)$
 - $F = G * h * f = G * (h * f)$

Mean filtering

- Find the average

$$\begin{array}{|c|c|c|} \hline 1/9 & 1/9 & 1/9 \\ \hline 1/9 & 1/9 & 1/9 \\ \hline 1/9 & 1/9 & 1/9 \\ \hline \end{array} * \begin{array}{|c|c|c|c|c|c|c|c|c|c|c|} \hline 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ \hline 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ \hline 0 & 0 & 0 & 90 & 90 & 90 & 90 & 90 & 0 & 0 & 0 \\ \hline 0 & 0 & 0 & 90 & 90 & 90 & 90 & 90 & 0 & 0 & 0 \\ \hline 0 & 0 & 0 & 90 & 90 & 90 & 90 & 90 & 0 & 0 & 0 \\ \hline 0 & 0 & 0 & 90 & 90 & 90 & 90 & 90 & 0 & 0 & 0 \\ \hline 0 & 0 & 0 & 90 & 0 & 90 & 90 & 90 & 0 & 0 & 0 \\ \hline 0 & 0 & 0 & 90 & 90 & 90 & 90 & 90 & 0 & 0 & 0 \\ \hline 0 & 0 & 0 & 90 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ \hline 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ \hline \end{array} = \begin{array}{|c|c|c|c|c|c|c|c|c|c|c|} \hline 0 & 10 & 20 & 30 & 30 & 30 & 20 & 10 & & & \\ \hline 0 & 20 & 40 & 60 & 60 & 60 & 40 & 20 & & & \\ \hline 0 & 30 & 60 & 90 & 90 & 90 & 60 & 30 & & & \\ \hline 0 & 30 & 50 & 80 & 80 & 90 & 60 & 30 & & & \\ \hline 0 & 30 & 50 & 80 & 80 & 90 & 60 & 30 & & & \\ \hline 0 & 20 & 30 & 50 & 50 & 60 & 40 & 20 & & & \\ \hline 10 & 20 & 30 & 30 & 30 & 30 & 20 & 10 & & & \\ \hline 10 & 10 & 10 & 0 & 0 & 0 & 0 & 0 & & & \\ \hline \end{array}$$

H F G

Linear filters: examples

- Identity filter:

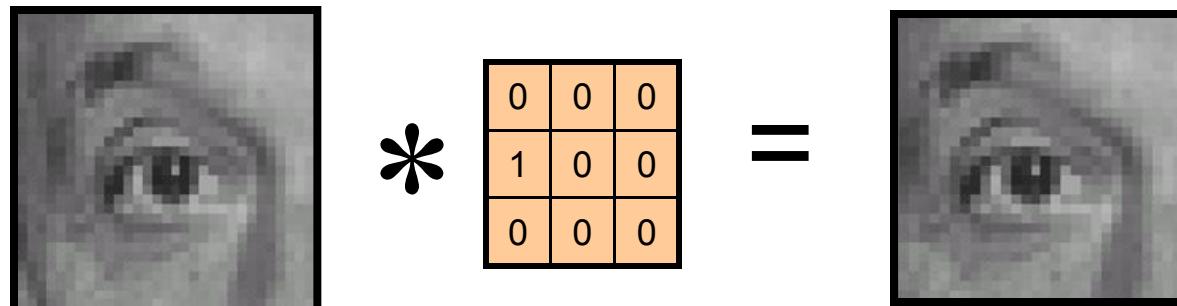
The diagram illustrates the application of a linear filter, specifically an identity filter, to a grayscale image of an eye. On the left, labeled "Original", is a grayscale image of an eye. In the center, a multiplication symbol (*) is positioned above a 3x3 matrix. The matrix has a value of 1 in the middle cell (row 2, column 2) and 0s in all other cells. To the right of the matrix is an equals sign (=). On the far right, labeled "Identical image", is another grayscale image of an eye, which appears identical to the original.

$$\text{Original} \quad * \quad \begin{matrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{matrix} = \text{Identical image}$$

Source: D. Lowe

Linear filters: examples

□ Shift Filter


$$\text{Original} \quad * \quad \begin{array}{|c|c|c|} \hline 0 & 0 & 0 \\ \hline 1 & 0 & 0 \\ \hline 0 & 0 & 0 \\ \hline \end{array} = \text{Shifted left By 1 pixel}$$

Source: D. Lowe

Linear filters: examples

□ Average/Blurring

The diagram illustrates the process of blurring an image using a mean filter. It consists of three main components: an original grayscale image of a face, a mathematical expression showing the convolution operation, and the resulting blurred image.

The original image is labeled "Original".

The mathematical expression shows the convolution operation: $\text{Original} * \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} = \text{Blur (with a mean filter)}$. The filter is a 3x3 matrix where every element is 1, and the divisor is 9.

The result of the convolution is a blurred grayscale image labeled "Blur (with a mean filter)".

Source: D. Lowe

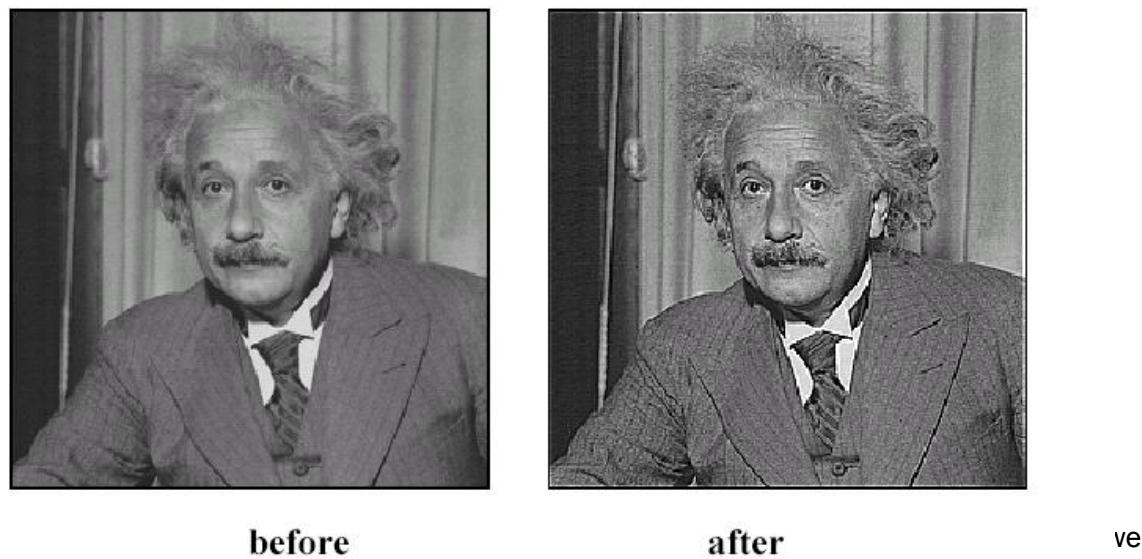
Linear filters: examples

- Sharpen: remove the average.

Original

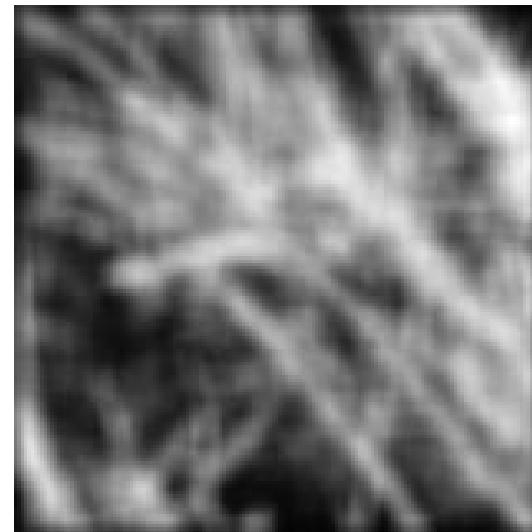
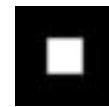
Sharpening filter

$$\text{Original} * \left(\begin{array}{ccc} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{array} - \frac{1}{9} \begin{array}{ccc} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{array} \right) = \text{Sharpened Image}$$



Smoothing with box filter

□ Box Filter:



□ Box filter in integral image domain:

- much faster, just 4 ADD operations.

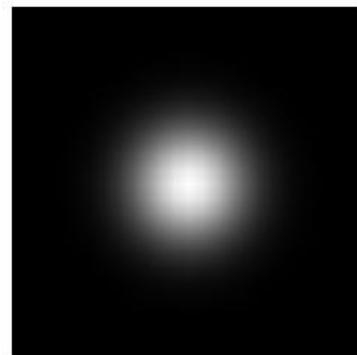
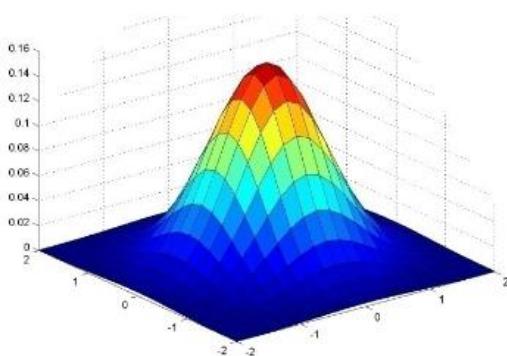
Source: D. Forsyth

Gaussian Kernel

- Gaussian Kernel with scale σ

$$G_\sigma = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

Matlab: `h=fspecial('gaussian', 5, 1.0);`

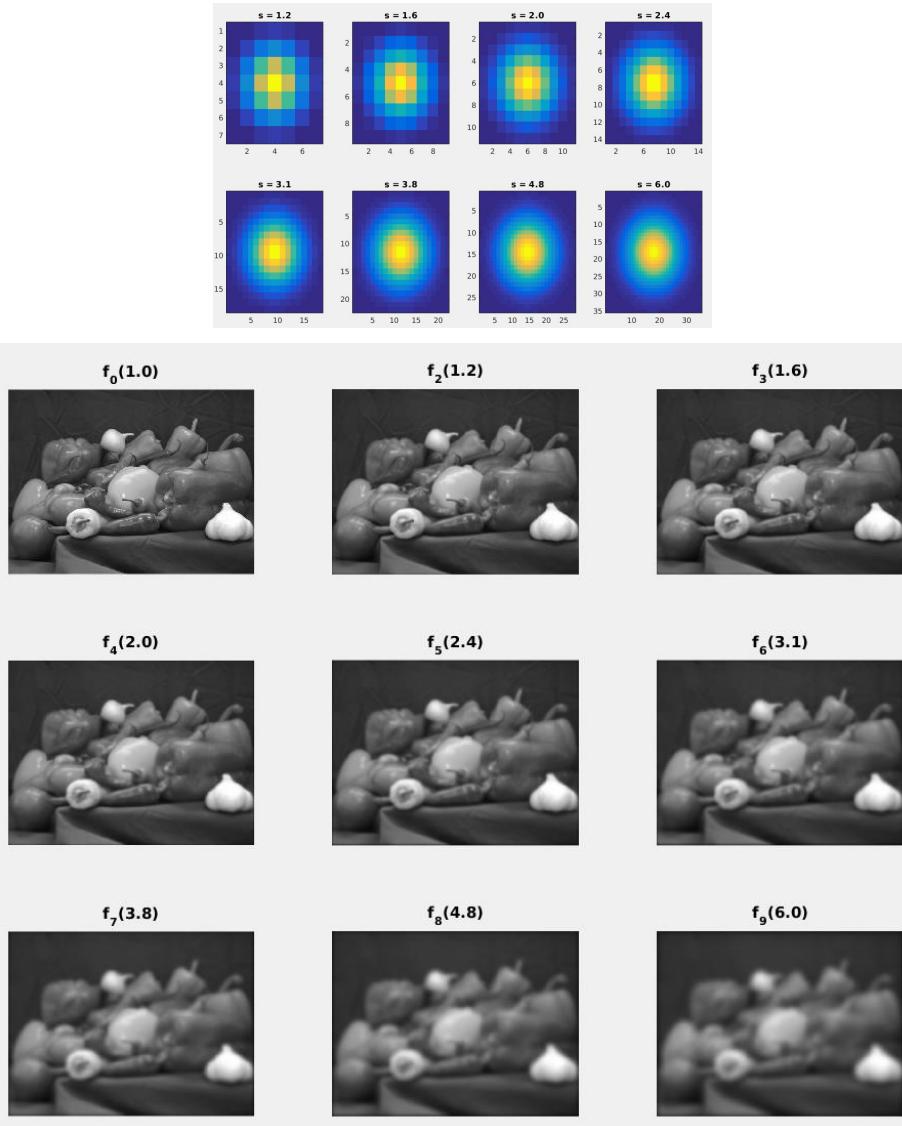


0.003	0.013	0.022	0.013	0.003
0.013	0.059	0.097	0.059	0.013
0.022	0.097	0.159	0.097	0.022
0.013	0.059	0.097	0.059	0.013
0.003	0.013	0.022	0.013	0.003

$5 \times 5, \sigma = 1$

Gaussian Filters – A Scale Space Approximation

□ Gaussian Blur



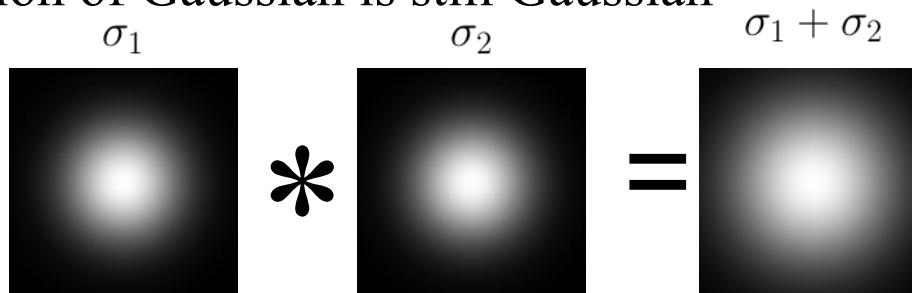
Matlab:

```
n=8;s = 1.25.^[1:n]; m=fix(6.*s);  
figure(30);  
for k=1:8  
    subplot(2,4,k); h = fspecial('gaussian',  
m(k), s(k));  
    imagesc(h); title(sprintf('s = %1.1f',  
s(k)));  
end  
  
figure(31);  
subplot(3,3,1); imshow(im);  
title('f_0(1.0)');  
for k=1:n  
    subplot(3,3,k+1);  
    h = fspecial('gaussian', m(k), s(k));  
    f = imfilter(im, h);  
    imshow(f);  
    title(sprintf('f_%d(%1.1f)', k+1, s(k)));  
end
```

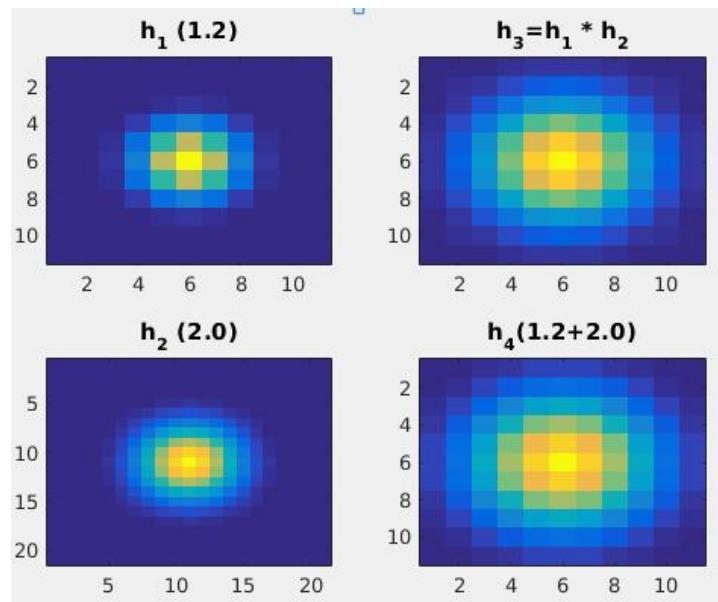
Gaussian filter properties

□ Successive Gaussian filtering

- convolution of Gaussian is still Gaussian

$$\sigma_1 \quad \quad \quad \sigma_2 \quad \quad \quad \sigma_1 + \sigma_2$$


new kernel sigma: $\text{h}_1 + \text{h}_2$

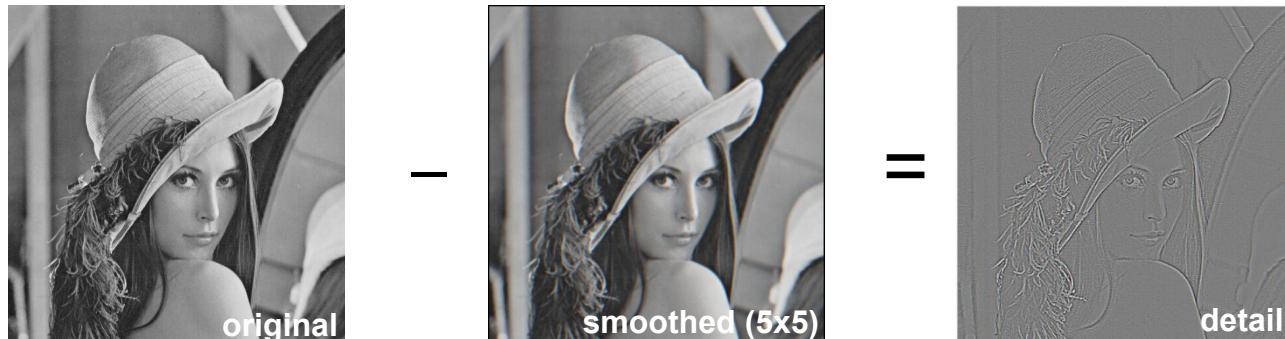


Matlab:

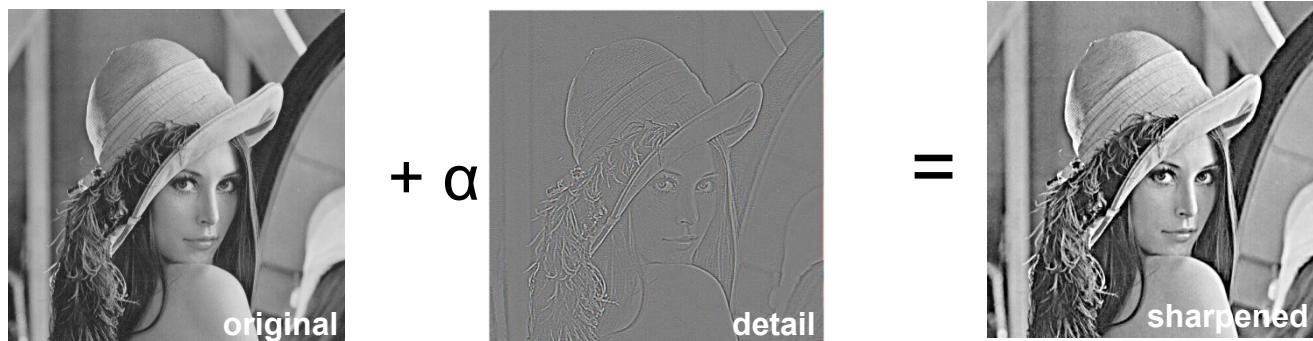
```
h1=fspecial('gaussian', 11, 1.2);
h2=fspecial('gaussian', 11, 2.0);
h3 = conv2(h1, h2);
h4=fspecial('gaussian', 11, 2.0+1.2);
```

Sharpening revisited

- What does blurring take away?



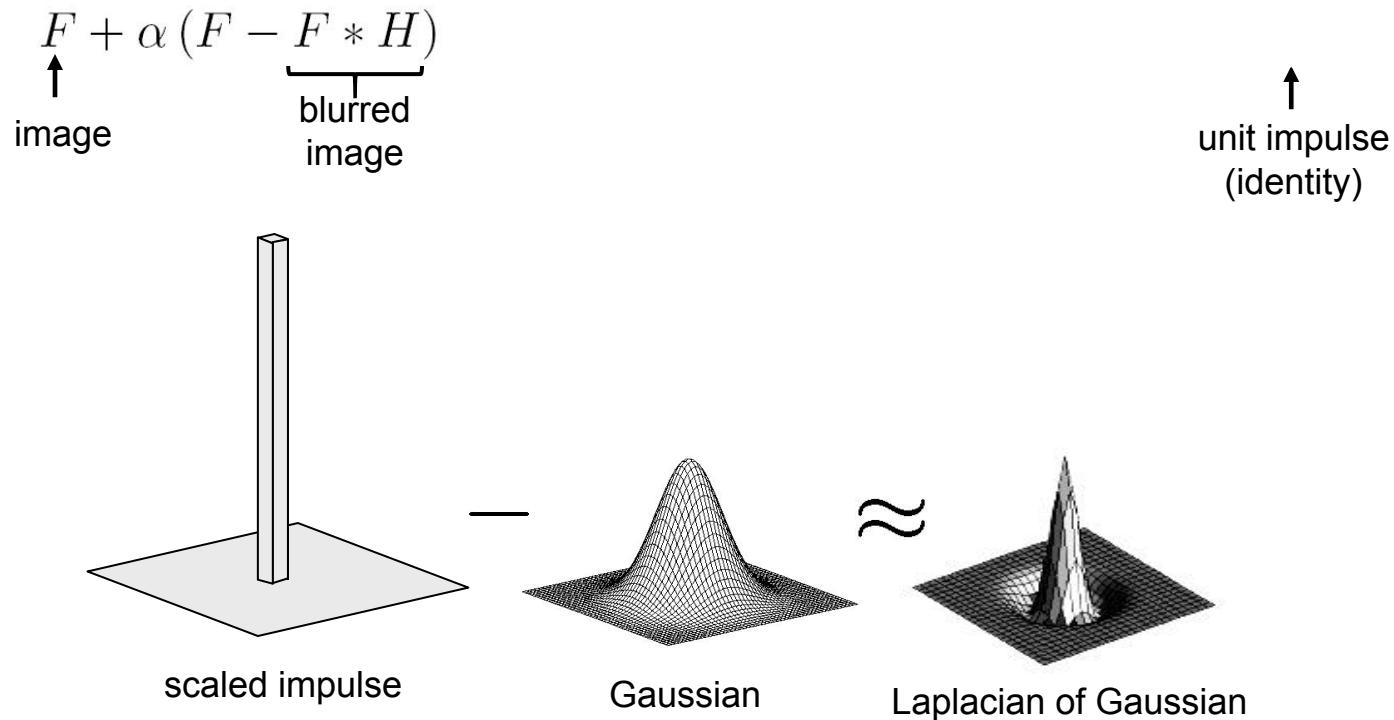
Let's add it back:



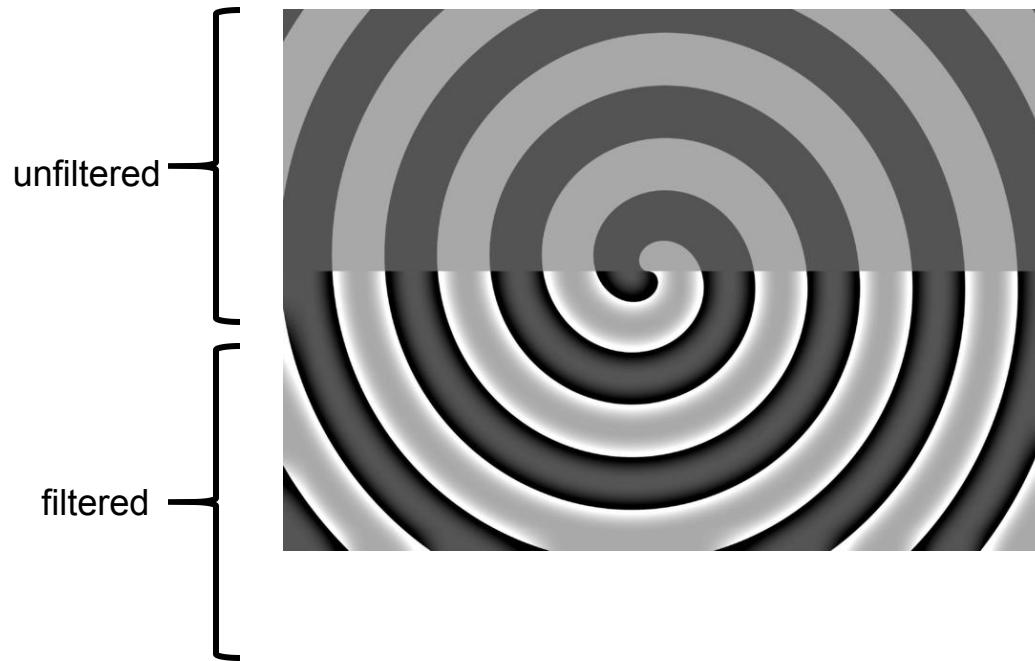
Source: S. Lazebnik

Sharpen filter - LoG

□ Laplacian of Gaussian

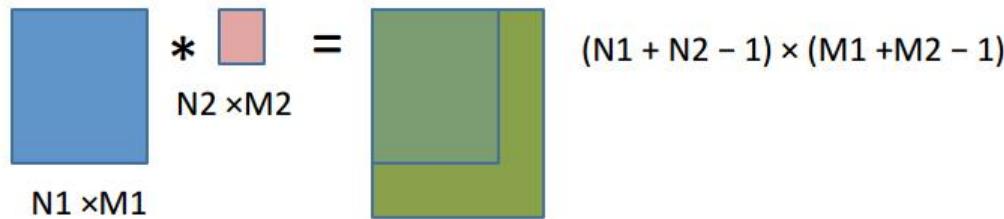


Sharpen filter



Filtering in Matlab

- Area of support for the operations



- To give the same $n_1 \times m_1$ output, need to padding the edge
 - Default is zero padding
 - Also replicate the last edge pixel
 - Or, mirroring (used in MPEG codec)

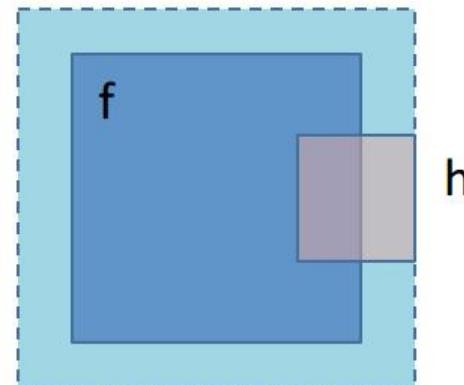
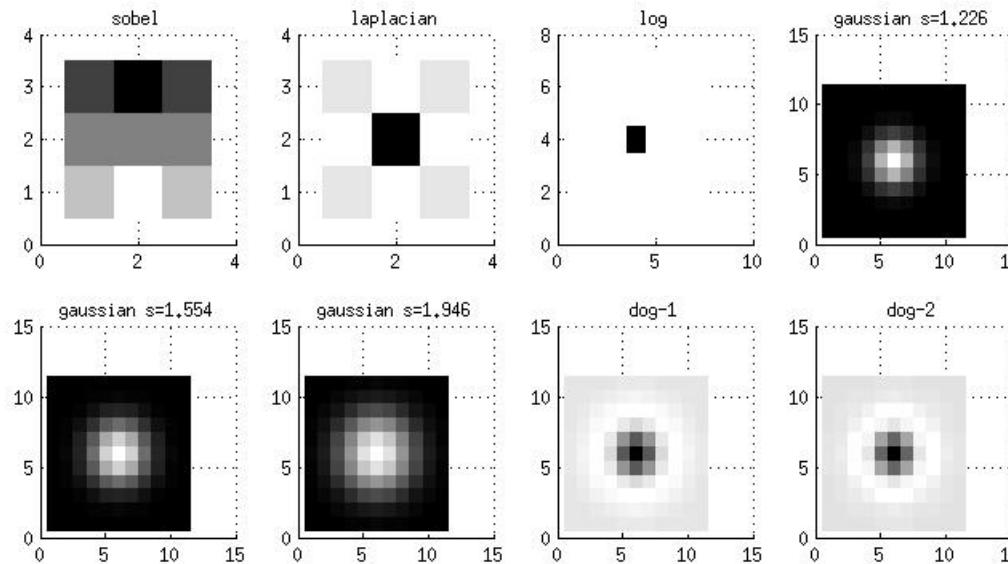


Image Filtering, Sweet Deal with Matlab

- It is such a nice tool
 - Main filter operation: `im2 = imfilter(im, h, 'replicate')`
 - Design your filter: `h= fspecial('filter_type', kernel_size, options)`
- Filter Design Examples:
 - Sobel
 - Laplacian, Laplacian of Gaussian
 - Gaussian, Difference of Gaussian (SIFT)



Matlab Image Filtering Example

```
%%%%%%%%%%%%%%
```

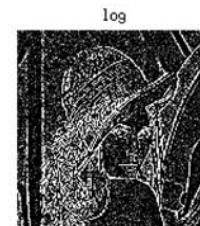
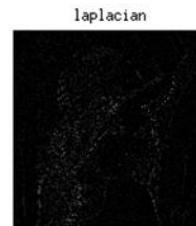
```
% image filters
```

```
%%%%%%%%%%%%%%
```

```
im = imread('..//pics/Lenna.png');
```

```
im = rgb2gray(im);
```

```
im1 = double(im(201:320, 201:320));
```



```
% edge filters
```

```
h{1} = fspecial('sobel');
```

```
h{2} = fspecial('laplacian', 0.25);
```

```
h{3} = fspecial('log', 7, 0.25);
```

```
% gaussians
```

```
sigmas = [1.226, 1.554, 1.946];
```

```
for k=1:length(sigmas)
```

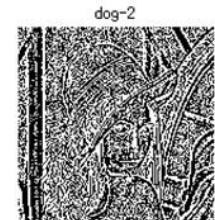
```
    h{3+k} = fspecial('gaussian',11, sigmas(k));
```

```
end
```

```
% diff of gaussian
```

```
h{7} = (h{6} - h{4}); h{7} = h{7}/sum(sum(h{7}));
```

```
h{8} = h{5} - h{4}; h{8} = h{8}/sum(sum(h{8}));
```



```
for k=1:8
```

```
    fprintf('\n k=%d', k);
```

```
    figure(26); subplot(2,4,k); grid on; hold on; colormap('gray'); imagesc(h{k});
```

```
    figure(27); subplot(2,4,k); imshow(imfilter(im, h{k}, 'replicate'));
```

```
end
```

Outline

- Recap of Lec 04
- Linear Filters
- Smoothing and Edge Detection
- Accelerating Linear Filters
- Summary

Filtering Properties

□ Linear operation that is

- Shift-Invariant:

$$f(m-k, n-j) * h = g(m-k, n-j), \text{ if } f * h = g$$

- Associative:

$$f * h_1 * h_2 = f * (h_1 * h_2)$$

this can save a lot of complexity

- Distributive:

$$f * h_1 + f * h_2 = f * (h_1 + h_2)$$

useful in SIFT's DoG filtering.

Separability of the Gaussian filter

- Gaussian smoothing is separable

$$\begin{aligned} G_\sigma(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

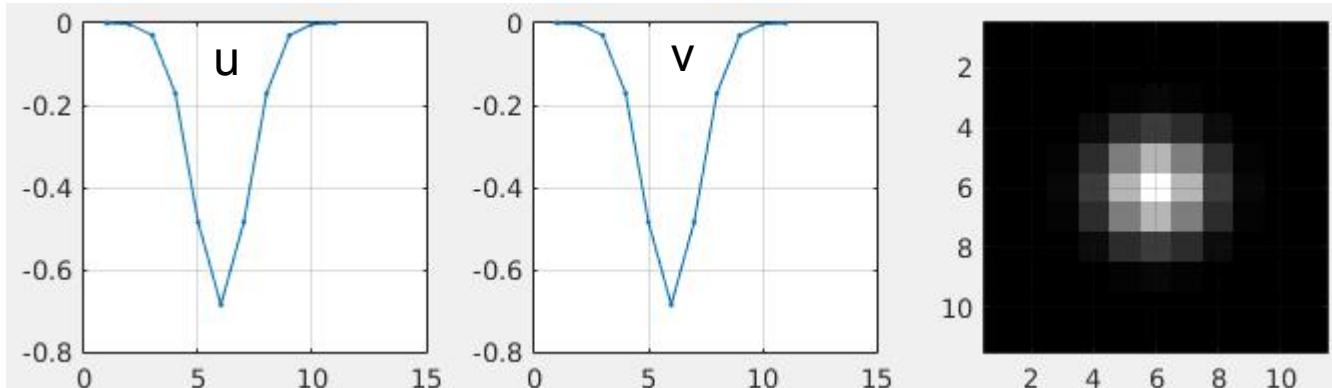
The 2D Gaussian can be expressed as the product of two functions, one a function of x and the other a function of y

In this case, the two functions are the (identical) 1D Gaussian

Separable Filter

- Gaussian filter is separable:

$$h = u \otimes v$$



- verify via SVD:
 - $h = \text{fspecial}('gaussian', 11, 1.2);$
 - $[u,s,v]=\text{svd}(h); \text{plot}(\text{diag}(s));$
- Benefits of separable filters: much faster
 - Implement as 1D filter by row with u , then transpose filter again by v .
 - instead of $O(m*n*k*k)$, $O(m*n*k)$

Separability example

□ Separable filtering complexity

2D convolution
(center location only)

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} * \begin{bmatrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{bmatrix}$$

The filter factors
into a product of 1D
filters:

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \times \begin{bmatrix} 1 & 2 & 1 \end{bmatrix}$$

Perform convolution
along rows:

$$\begin{bmatrix} 1 & 2 & 1 \end{bmatrix} * \begin{bmatrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{bmatrix} = \begin{bmatrix} 11 \\ 18 \\ 18 \end{bmatrix}$$

Followed by convolution
along the remaining column:

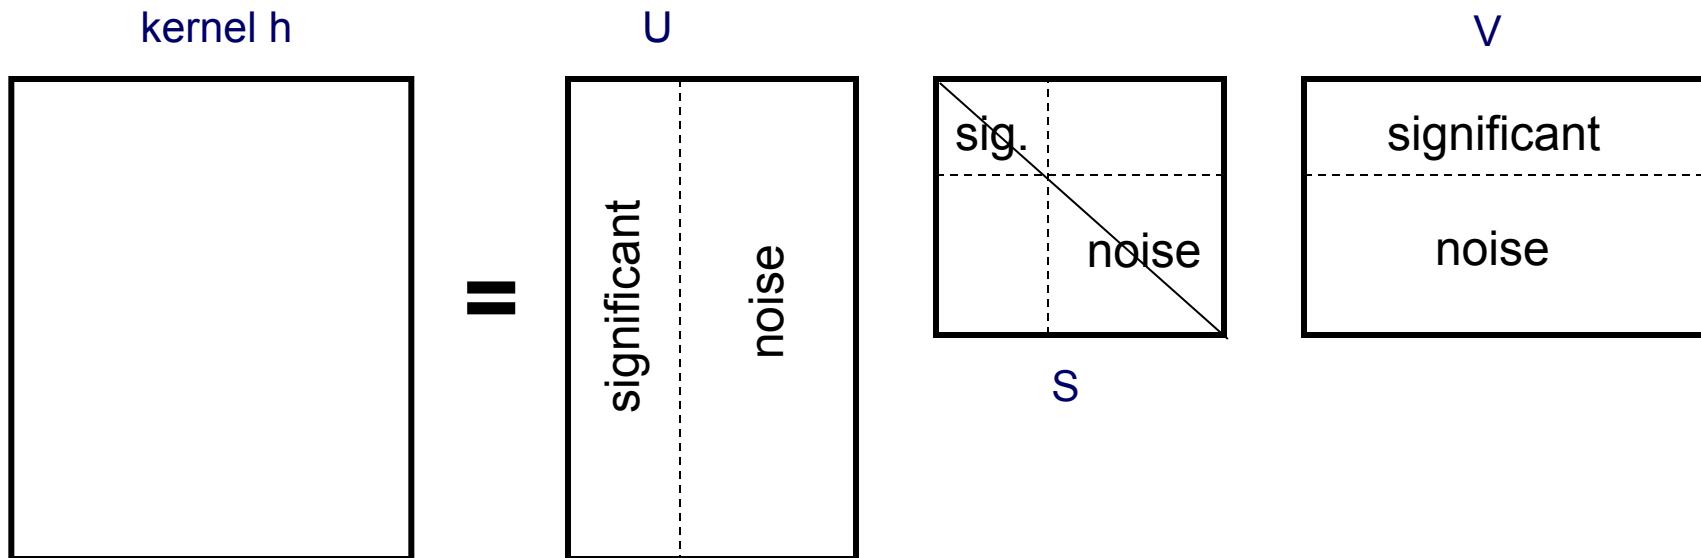
$$\begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} * \begin{bmatrix} 11 \\ 18 \\ 18 \end{bmatrix} = \begin{bmatrix} 65 \end{bmatrix}$$

Approx. Non-Separable Filters

- ❑ k-SVD separable filters
 - filter kernel is approximated by first SVD basis

$$[U, S, V] = svd(h)$$

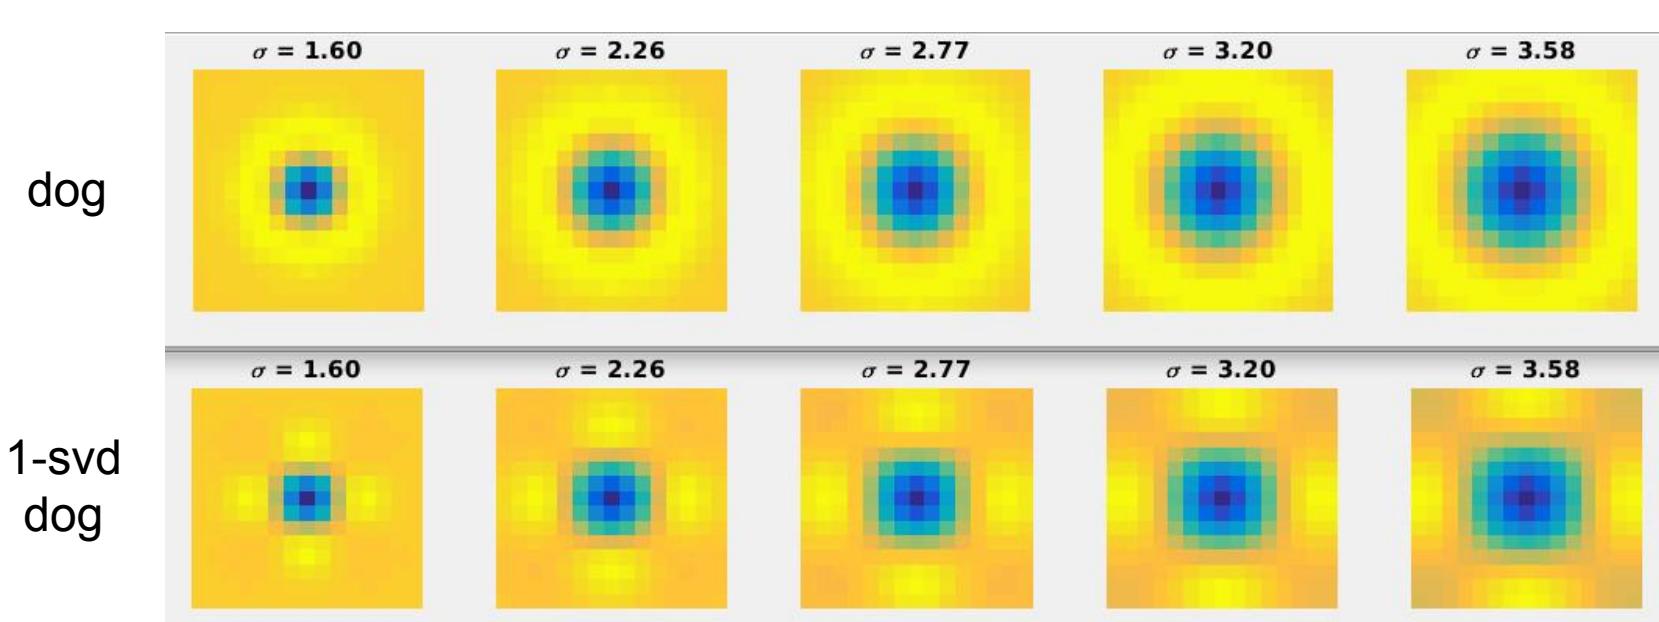
$$h = \sum_k s_{k,k} u_k v_k T$$



Matlab Implementation

□ DoG approximation with k-SVD separable filtering

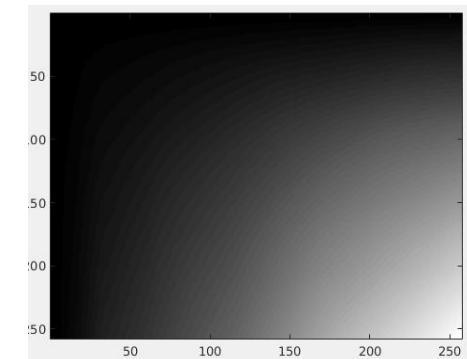
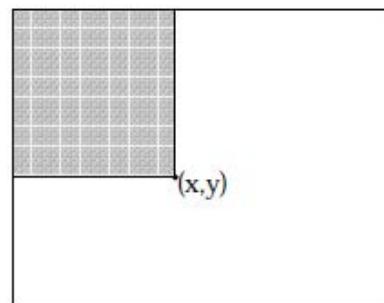
```
% k-SVD DoG filter
for k=1:length(sigma) -1
    [u,s,v]=svd(h_dog{k});
    % 1-SVD approx
    dog_svd{k} = s(1,1)*u(:,1)*v(:,1)';
    ss(k,:) = diag(s);
    figure(24); subplot(1,5,k); imagesc(h_dog{k});
    figure(25); subplot(1,5,k); imagesc(dog_svd{k});
end
```



2D Box Filtering

❑ Integral Images:

$$ii(x, y) = \sum_{x' \leq x, y' \leq y} i(x', y')$$



❑ Fast 2D box filtering

- applying a rectangle/box filter can be computed as sum of 4 corners
- much faster than convolution

❑ One of the key applications:

- AdaBoost face detection

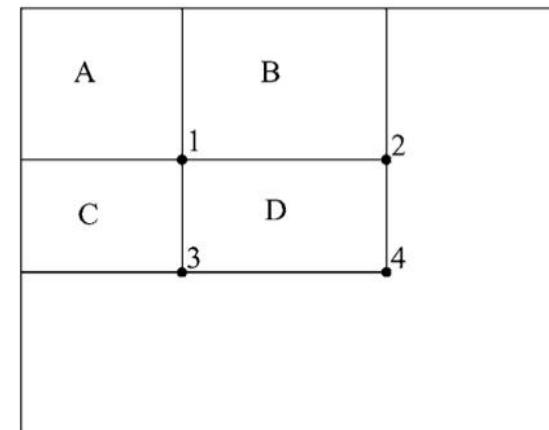
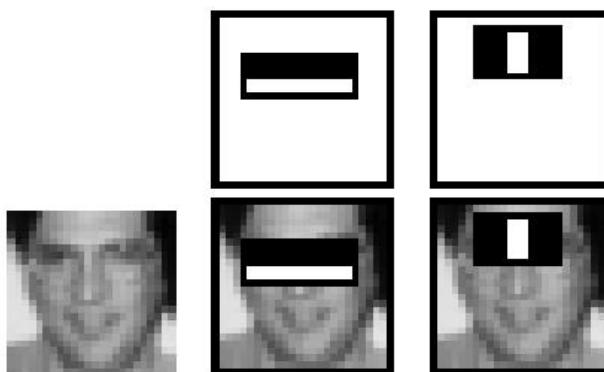


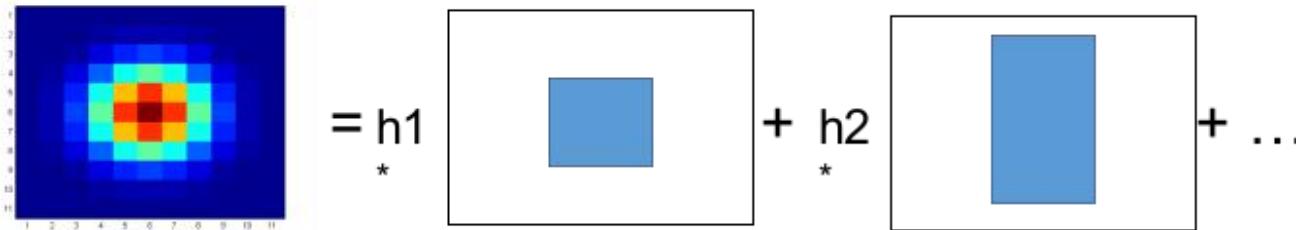
Figure 3. The sum of the pixels within rectangle D can be computed with four array references. The value of the integral image at location 1 is the sum of the pixels in rectangle A . The value at location 2 is $A + B$, at location 3 is $A + C$, and at location 4 is $A + B + C + D$. The sum within D can be computed as $4 + 1 - (2 + 3)$.

Complexity Reduction

- The overhead of a 2D convolution can be significant, especially when the kernel is big.
 - Cost: $O(MNPQ)$ multiplications and $O(MNPQ)$ additions for an image of size MxN and a kernel of size PxQ.
- The complexity can be alleviated if the kernel is separable.
 - Cost: $O(MN(P + Q))$ multiplications and $O(MN(P + Q))$ additions.
- For large images and kernels the convolution can still be computationally expensive, e.g., detecting key points using Difference of Gaussians (DoG.)
- What if we could use Integral images and box filters to approximate the original filter?
 - The cost of computing an Integral image is $O(MN)$ additions
 - The cost of convolving with a k-box linear combination filters approximation is $O(MNK)$ multiplications and $O(4MNK)$ additions, where k is the number of boxes used for the approximation.

Approximating Gaussian Filter

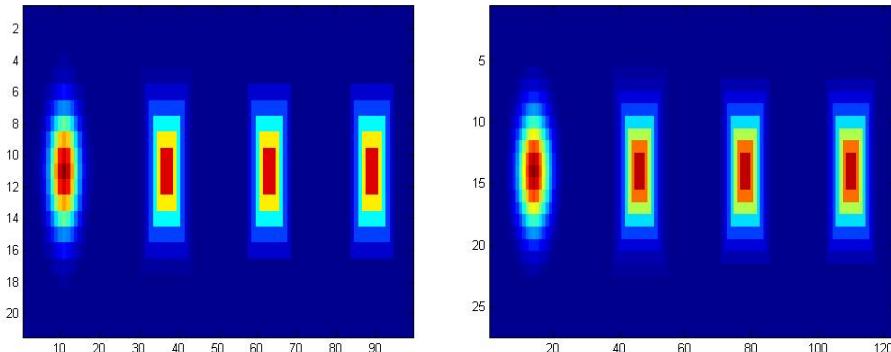
- Idea: A filter can be approximated as a linear combination of several box filters.



We need to find the set of weights $\{h_i\}$, such that:

- the number of boxes to use is as minimal as possible
- keep a low approximation error

Sigma	Residuals			Num. of Boxes		
	LS	LS-LASSO	LASSO	LS	LS-LASSO	LASSO
1.249	0.05535	0.055372	0.055372	5	3	3
1.226273	0.057776	0.057791	0.057791	5	3	3
1.545008	0.035789	0.035802	0.035802	7	4	4
1.946588	0.024847	0.024848	0.024848	8	6	6
2.452547	0.018447	0.019152	0.019152	10	5	5
3.090016	0.014137	0.014175	0.014175	13	8	8



Application: CABOX detection

- We use the dictionary generation approach (3) and LASSO.



- The box filter (red circles) approach was able to detect **88%** of the features detected by the Gaussian approach (green circles).
- Another experiment with 190 images revealed that the box filter approach was able to detect **65%** of the features detected by Vlfeat.
- The criterion for declaring overlapping features is if their distance is less than 5 pixels (typically used in SfM)

Summary

□ Linear Filters

- involves a neighbourhood in output pixel values computing
- characterized by a kernel: $y = \text{conv2}(x, h)$
- **Shift-Invariant:**
 $f(m-k, n-j)*h = g(m-k, n-j)$, if $f*h=g$
- **Associative:**
 $f*h_1*h_2 = f*(h_1*h_2)$
this can save a lot of complexity
- **Distributive:**
 $f*h1 + f*h2 = f*(h1+h2)$ useful in SIFT's DoG filtering.

□ Smoothing/Edge Detection

- Gaussian blurring
- DoG edge detection

□ Filtering acceleration

- Separable filter
- Integral image domain box filtering